

Network Traffic Classification Using Machine Learning and Deep Learning Approaches






















Liu, Yucheng



Research Background

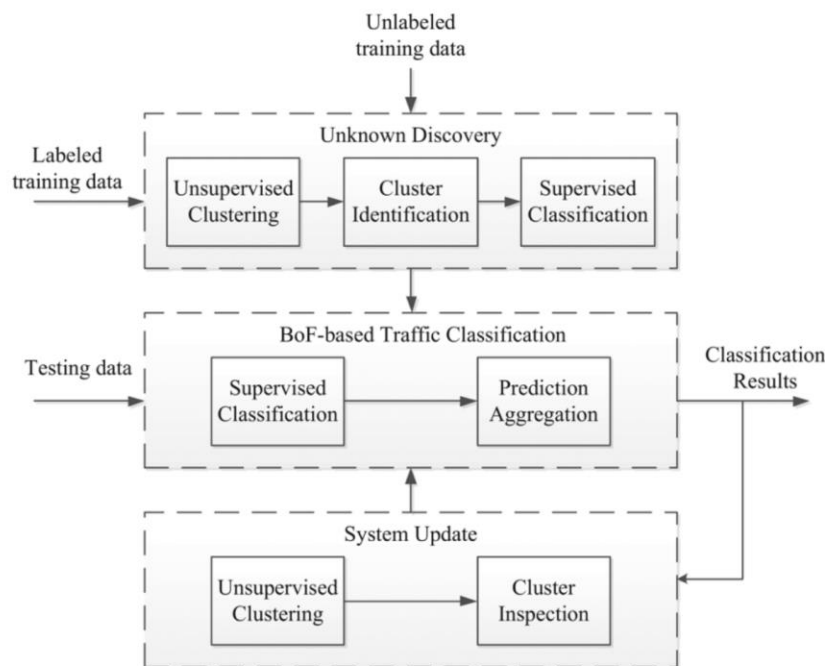
Project Objectives

- **Flow classification goal:** map raw encrypted traffic → application labels (e.g. Chat vs. Video vs. P2P)
- **End-to-end pipeline:** PCAP → feature extraction → ML/DL model → real-time inference
- **Key questions:**
- Can lightweight ML (K-NN, RF) match deep models?
- Does temporal modeling (CNN+RNN) improve accuracy?
- What is the training time vs. accuracy trade-off?

 AUDIO_spotifygateway.pcap
 BROWSING_gate_SSL_Browsing.pcap
 BROWSING_ssl_browsing_gateway.pcap
 CHAT_aimchatgateway.pcap
 CHAT_facebookchatgateway.pcap
 CHAT_gate_AIM_chat.pcap
 CHAT_gate_facebook_chat.pcap
 CHAT_gate_hangout_chat.pcap
 CHAT_gate_ICQ_chat.pcap
 CHAT_gate_skype_chat.pcap
 CHAT_hangoutschatgateway.pcap
 CHAT_icqchatgateway.pcap
 CHAT_skypechatgateway.pcap
 FILE-TRANSFER_gate_FTP_transfer.pcap
 FILE-TRANSFER_gate_SFTP_filetransfer.pcap
 MAIL_gate_Email_IMAP_filetransfer.pcap
 MAIL_gate_POP_filetransfer.pcap
 MAIL_Gateway_Thunderbird_Imap.pcap
 MAIL_Gateway_Thunderbird_POP.pcap
 torFacebook.pcap
 torGoogle.pcap


Project Objectives

- **Flow classification goal:** map raw application labels (e.g. Chat vs. Video)
- **End-to-end pipeline:** PCAP → feature extraction → model → real-time inference
- **Key questions:**
- Can lightweight ML (K-NN, RF) match accuracy of deep learning?
- Does temporal modeling (CNN+RNN) improve accuracy?
- What is the training time vs. accuracy trade-off?



Project Objectives

- **Flow classification goal:** map raw encrypted traffic → application labels (e.g. Chat vs. Video vs. P2P)
- **End-to-end pipeline:** PCAP → feature extraction → ML/DL model → real-time inference
- **Key questions:**
 - Can lightweight ML (K-NN, RF) match deep models?
 - Does temporal modeling (CNN+RNN) improve accuracy?
 - What is the training time vs. accuracy trade-off?




MACHINE LEARNING VS DEEP LEARNING

AUDIO_spotifygateway.pcap
BROWSING_gate_SSL_Browsing.pcap
BROWSING_ssl_browsing_gateway.pcap

MAIL_Gateway_Thunderbird_Imap.pcap
MAIL_Gateway_Thunderbird_POP.pcap
torFacebook.pcap
torGoogle.pcap

Project Objectives

- **Flow classification goal:** map raw encrypted traffic → application labels (e.g. Chat vs. Video vs. P2P)
- **End-to-end pipeline:** PCAP → feature extraction → ML/DL model → real-time inference
- **Key questions:**
 - Can lightweight ML (K-NN, RF) match deep models?
 - Does temporal modeling (CNN+RNN) improve accuracy?
 - What is the training time vs. accuracy trade-off?



Key Challenges

- **Raw PCAP complexity:** packet parsing, flow reassembly, noise removal
- **Feature selection:** hundreds of statistical metrics; choosing meaningful subset
- **Capturing temporal patterns:** per-flow aggregates lose intra-flow dynamics / need sub-window sequence representation

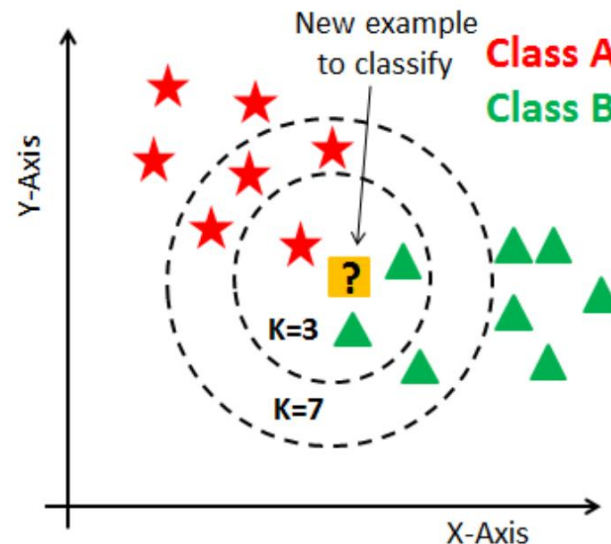
Key Challenges

- **Raw PCAP complexity:** packet parsing, flow reassembly, noise removal
- **Feature selection:** hundreds of statistical metrics; choosing meaningful subset
- **Capturing temporal patterns:** per-flow aggregates lose intra-flow

```
Source IP, Source Port, Destination IP, Destination Port, Protocol, Flow Duration,  
Flow Bytes/s, Flow Packets/s, Flow IAT Mean, Flow IAT Std, Flow IAT Max, Flow IAT Min,  
Fwd IAT Mean, Fwd IAT Std, Fwd IAT Max, Fwd IAT Min, Bwd IAT Mean, Bwd IAT Std,  
Bwd IAT Max, Bwd IAT Min, Active Mean, Active Std, Active Max, Active Min, Idle Mean,  
Idle Std, Idle Max, Idle Min, label
```


Review – Traditional ML

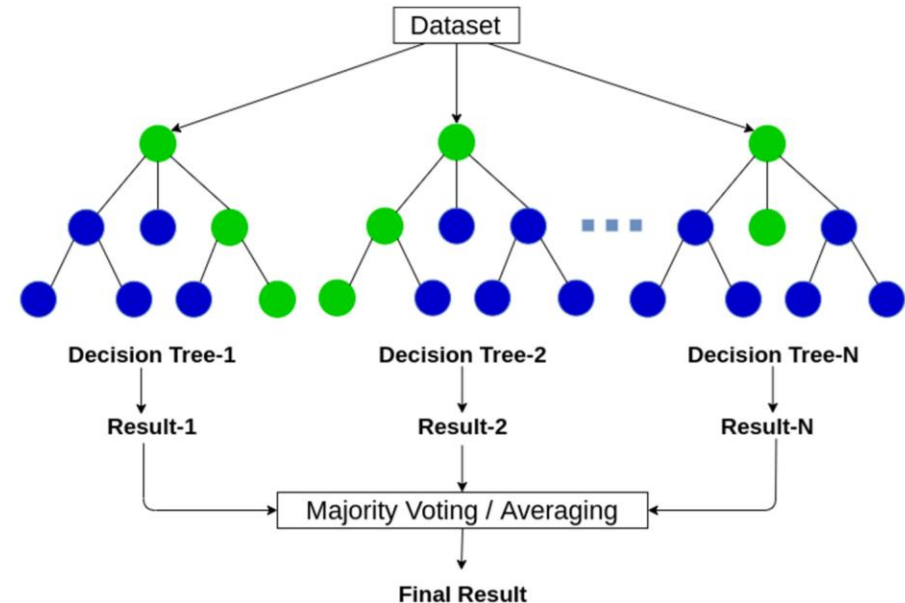
- **K-Nearest Neighbors (K-NN)**
 - simple, non-parametric; labels by nearest feature vectors
 - *Pros*: no training phase, interpretable
 - *Cons*: slow at inference, sensitive to feature scaling



Review – Traditional ML

- **Random Forest (RF)**

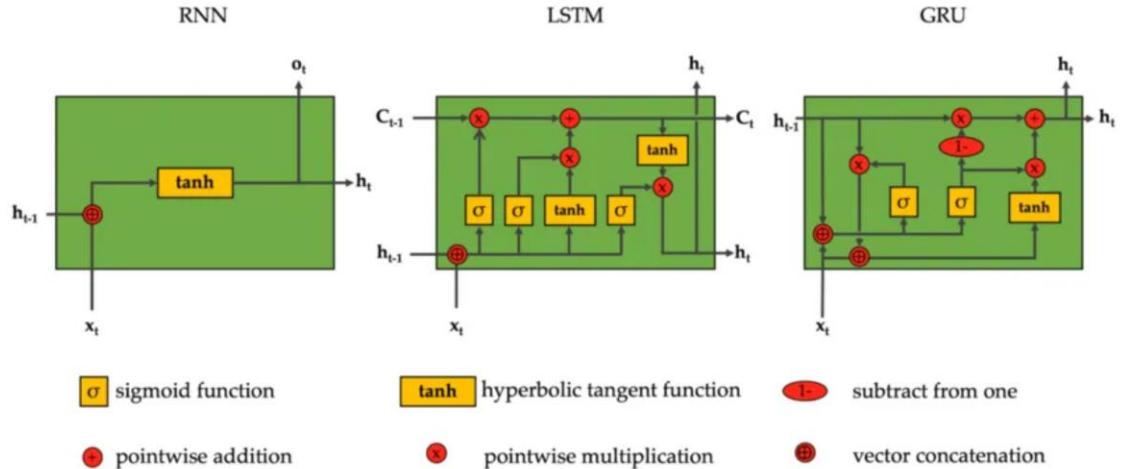
- Ensemble of decision trees trained on random subsets of samples & features
- **Pros:** Robust to noise and overfitting, Handles high-dimensional data well
- **Cons:** Large model size, Longer training time as number of trees grows



Review – Deep Learning Method

- RNN family (LSTM / GRU)**

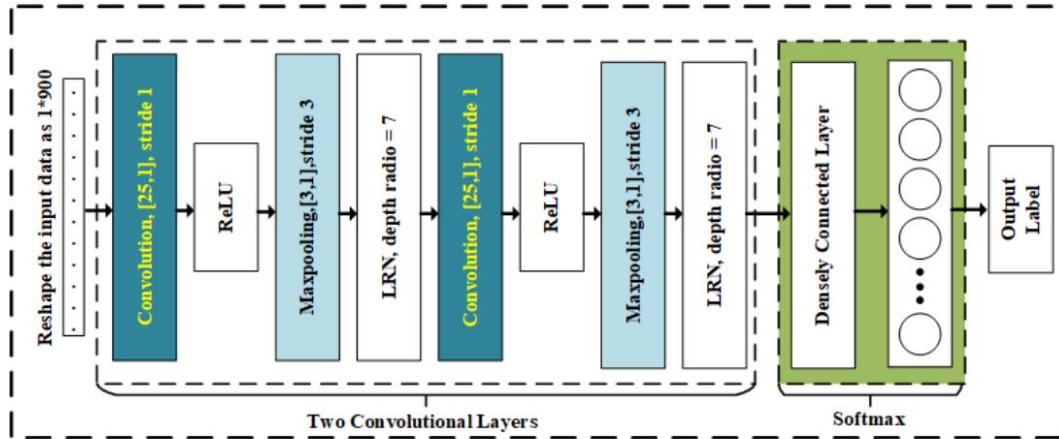
- model long/short-term temporal dependencies across windows
- handle variable-length flows, resistant to jitter



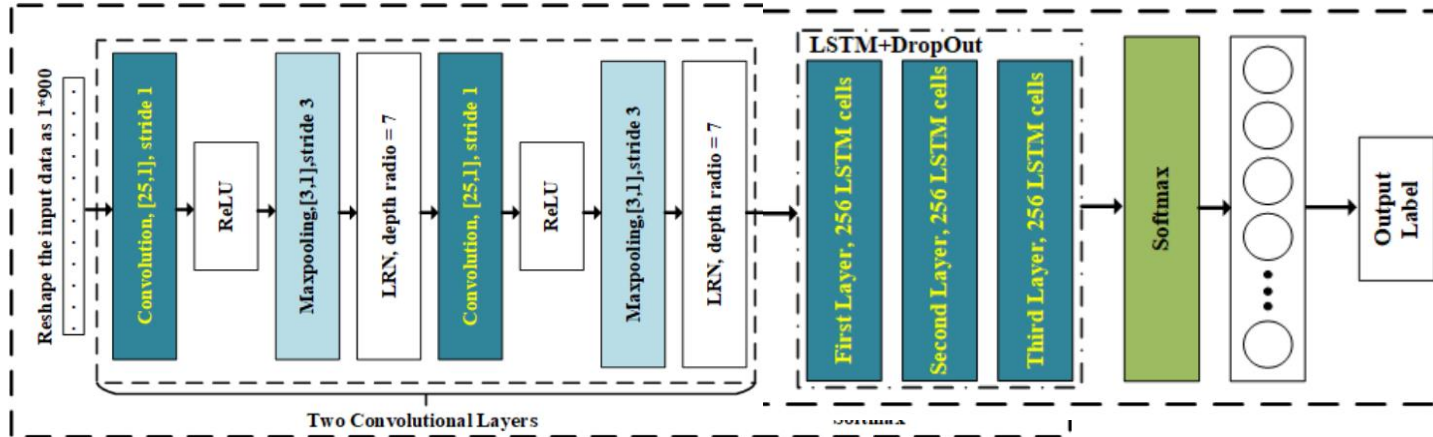
Review – Deep Learning Method

- **Hybrid architectures**
 - **CNN + LSTM/GRU** pipeline (Zeng et al., IEEE Access 2019 “Deep-Full-Range”)
 - integrate spatial & temporal modules for best of both worlds


Review – Deep Learning Method



Review – Deep Learning Method



Review – Deep Learning Method



Deep-Full-Range: A Deep Learning Based Network Encrypted Traffic Classification and Intrusion Detection Framework

Yi Zeng; Huaxi Gu

Algorithm 5 Online-Fashioned DFR Algorithm

Input:

$RT^1, RT^2, \dots, RT^t, \dots, RT^T$

Output:

The DFR model that best fits with the current traffic environment

```

1: for each  $t$  in  $(1, T)$  do
2:   Apply Alg. 1 to gain  $G(1, 2, \dots, j, \dots, J)$ 
3:   Randomly separate  $G(1, 2, \dots, j, \dots, J)$  according to 9:1;
4:   Use  $G_{Train}$  to train the three DFRs;
5:   Use  $G_{Test}$  to select the highest-accuracy DFR;
6:   Run the current-using DFR on the same  $G_{Test}$ ;
7:   if the current-using DFR's accuracy is smaller then
8:     Save the new DFR;
9:     Transmit the new DFR to other units
10:  end if
11: end for
    
```

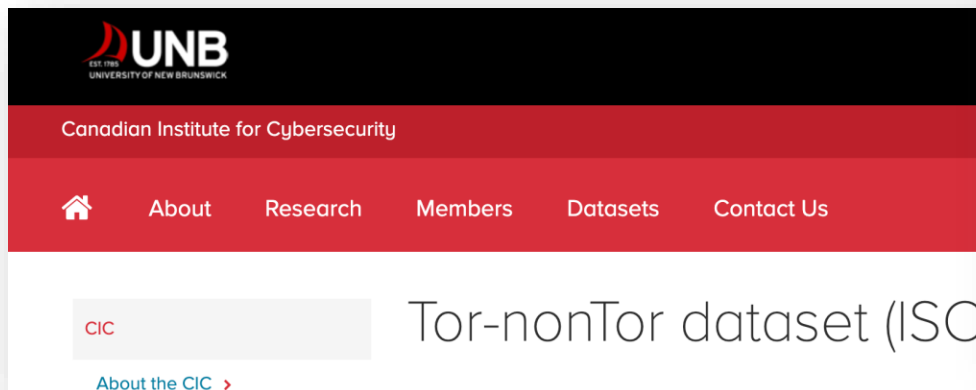
Preparation

`./pcap_analysis.py`

Data Processing & Feature Extraction

- **PCAP File Collection**

- Collect network traffic data in PCAP format
- Extract traffic labels from filenames



- AUDIO_spotifygateway.pcap
- BROWSING_gate_SSL_Browsing.pcap
- BROWSING_ssl_browsing_gateway.pcap
- CHAT_aimchatgateway.pcap
- CHAT_facebookchatgateway.pcap
- CHAT_gate_AIM_chat.pcap
- CHAT_gate_facebook_chat.pcap
- CHAT_gate_hangout_chat.pcap
- CHAT_gate_ICQ_chat.pcap
- CHAT_gate_skype_chat.pcap
- CHAT_hangoutschatgateway.pcap
- CHAT_icqchatgateway.pcap
- CHAT_skypechatgateway.pcap
- FILE-TRANSFER_gate_FTP_transfer.pcap
- FILE-TRANSFER_gate_SFTP_filetransfer.pcap
- MAIL_gate_Email_IMAP_filetransfer.pcap
- MAIL_gate_POP_filetransfer.pcap
- MAIL_Gateway_Thunderbird_Imap.pcap
- MAIL_Gateway_Thunderbird_POP.pcap
- torFacebook.pcap
- torGoogle.pcap

Data Processing

- Group packets by 5-tuple (src_ip, src_port, dst_ip, dst_port, protocol)
 - Split flows into 10-second intervals for consistent analysis
 - Separate forward and backward traffic directions

```
for packet in packets:
    if IP in packet:
        if TCP in packet:
            proto = 'TCP'
            src_port = packet[TCP].sport
            dst_port = packet[TCP].dport
```

```
# Separate forward and backward packets
fwd_packets = [p for p in packets if p[IP].src == src_ip]
bwd_packets = [p for p in packets if p[IP].src == dst_ip]

# Calculate forward IAT statistics
fwd_times = [p.time for p in fwd_packets]
fwd_iats = [fwd_times[i+1] - fwd_times[i] for i in range(len(fwd_times)-1)]
```

```
if src_ip < dst_ip:
    flow_key = (src_ip, src_port, dst_ip, dst_port, proto)
else:
    flow_key = (dst_ip, dst_port, src_ip, src_port, proto)

flows[flow_key].append(packet)

return flows
```

Feature Extraction – Basic Features

- Flow Duration, Flow Bytes
- IAT Statistics (Mean, Std,
- Overall Flow
- Forward Direction
- Backward Direction
- Active/Idle Time Statistics

```
# Calculate flow bytes and packets per second
```

```
total_bytes = sum(len(n) for n in packets)
```

```
# Calculate packet inter-arrival times (IAT)
```

```
packet_times = [p.time for p in packets]
```

```
iats = [packet_times[i+1] - packet_times[i] for i in range(len(packet_times)-1)]
```

```
# Calculate active and idle times
```

```
if iats[i] > idle_threshold:
```

```
    # End of active period
```

```
    active_time = packet_times[i] - current_active_start
```


```
    active_times.append(active_time)
```

```
    idle_times.append(iats[i])
```

```
    current_active_start = packet_times[i+1]
```

Feature Extraction - Temporal Features

- **Advanced Temporal Features (80 features)**
- 10 time windows × 8 features per window:
- Packets_Count, Bytes_Count
- IAT_Mean, IAT_Std
- Fwd_IAT_Mean, Fwd_IAT_Std
- Bwd_IAT_Mean, Bwd_IAT_Std



Data Processing

- **Preprocessing Steps:**
- Remove IP addresses and ports
- Handle missing values
- Replace infinite values
- Standardize features (zero mean, unit variance)
- Split data (80% training, 20% testing)

My Model Architecture

`./base_model.py`
`./advanced_model.py`

KNN

- **Hyperparameter Optimization:** GridSearchCV with 5-fold cross-validation
- **K values tested:** 3, 5, 7, 9, 11, 13, 15
- **Distance Metric:** Euclidean (default)
- **Weight Function:** Uniform weighting (default)

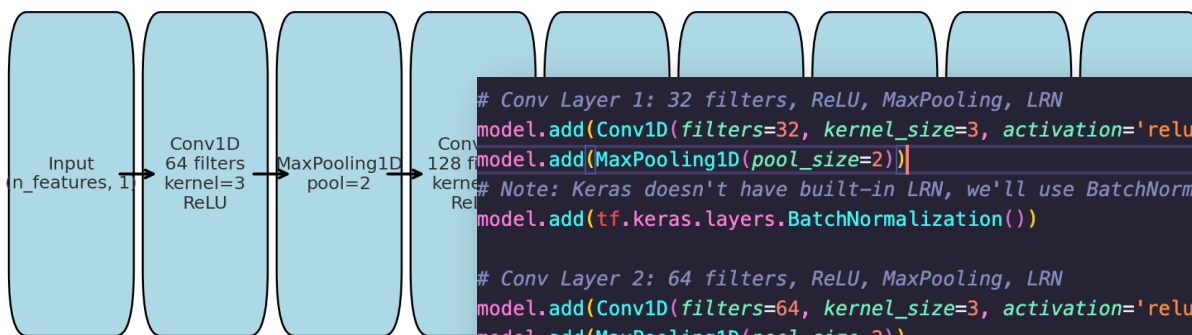
```
def train_knn(X_train, y_train, X_test, y_test, feature_names):  
    """Train and evaluate KNN model"""  
    print("Training KNN model...")  
  
    # Use grid search to find the best K value  
    param_grid = {'n_neighbors': [3, 5, 7, 9, 11, 13, 15]}  
    knn = KNeighborsClassifier()  
    grid_search = GridSearchCV(knn, param_grid, cv=5, scoring='accuracy')  
    grid_search.fit(X_train, y_train)  
  
    # Get the best model  
    best_knn = grid_search.best_estimator_  
    best_k = grid_search.best_params_['n_neighbors']  
    print(f"Best K value: {best_k}")
```

Random Forest

- Trees in Forest: 10, 50, or 100 (optimized via GridSearchCV)
- Max Tree Depth: None , 10, or 20
- Min Samples to Split: 2, 5, or 10
- Criterion: Gini impurity (default)
- Bootstrap: True (default)

```
def train_random_forest(X_train, y_train, X_test, y_test, feature_names):  
    """Train and evaluate Random Forest model"""  
    print("Training Random Forest model...")  
  
    # Define parameter grid for grid search  
    param_grid = {  
        'n_estimators': [10, 50, 100],  
        'max_depth': [None, 10, 20],  
        'min_samples_split': [2, 5, 10]  
    }  
  
    # Create Random Forest classifier  
    rf = RandomForestClassifier(random_state=42)  
  
    # Perform grid search with cross-validation  
    grid_search = GridSearchCV(rf, param_grid, cv=5, scoring='accuracy')  
    grid_search.fit(X_train, y_train)  
  
    # Get best model  
    best_rf = grid_search.best_estimator_  
    best_params = grid_search.best_params_  
    print(f"Best parameters: {best_params}")
```


1D-CNN



```
# Conv Layer 1: 32 filters, ReLU, MaxPooling, LRN
model.add(Conv1D(filters=32, kernel_size=3, activation='relu', input_shape=(X_train.shape[1], 1)))
model.add(MaxPooling1D(pool_size=2))
# Note: Keras doesn't have built-in LRN, we'll use BatchNormalization as an alternative
model.add(tf.keras.layers.BatchNormalization())

# Conv Layer 2: 64 filters, ReLU, MaxPooling, LRN
model.add(Conv1D(filters=64, kernel_size=3, activation='relu'))
model.add(MaxPooling1D(pool_size=2))
model.add(tf.keras.layers.BatchNormalization())

# Flatten and Dense layers
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5)) # Adding dropout for regularization
model.add(Dense(num_classes, activation='softmax'))
```

CNN + LSTM

CNN Layers:

Conv1D: 64 filters, kernel size 3, ReLU activation, same padding

MaxPooling1D: pool size 2

Conv1D: 128 filters, kernel size 3, ReLU activation, same padding

MaxPooling1D: pool size 2

RNN Layers:

Bidirectional LSTM: 128 units, return sequences=True

Dropout: 30%

LSTM: 64 units

Dropout: 30%

CNN + GRU (simplified version LSTM)

CNN Layers:

Conv1D: 64 filters, kernel size 3, ReLU activation, same padding

MaxPooling1D: pool size 2

Conv1D: 128 filters, kernel size 3, ReLU activation, same padding

MaxPooling1D: pool size 2

RNN Layers:

Bidirectional LSTM: 128 units, return sequences=True

Dropout: 30%

GRU: 64 units

Dropout: 30%

LSTM vs GRU

Internal Architecture

LSTM: Has 3 gates (input, forget, output)

GRU: Has 2 gates (update, reset)

Parameter Count

LSTM: 4 weight matrices (more parameters)

GRU: 3 weight matrices (fewer parameters)

CNN-LSTM Model:


```
# LSTM part
x = Bidirectional(LSTM(128, return_sequences=True))(x)
x = Dropout(0.3)(x)
x = LSTM(64)(x)
x = Dropout(0.3)(x)
```

CNN-GRU Model:


```
# GRU part
x = Bidirectional(GRU(128, return_sequences=True))(x)
x = Dropout(0.3)(x)
x = GRU(64)(x)
x = Dropout(0.3)(x)
```

Result Analysis


Accuracy




Time Spent



Network Traffic Classification Using ML & DL approaches



KNN




RF


Network Traffic Classification Using ML & DL approaches

KNN

KNN Model - Top 10 Important Features




RandomForest Model - Top 10 Important Features




RandomForest

Network Traffic Classification Using ML & DL approaches




Learning_Curve - KNN




Learning_Curve - RF


Network Traffic Classification Using ML & DL approaches



CNN_GRU






CNN_LSTM




1D_CNN

Network Traffic Classification Using ML & DL approaches





1D_CNN


Network Traffic Classification Using ML & DL approaches





CNN_LSTM




Network Traffic Classification Using ML & DL approaches




CNN_GRU



Precision Recall




KNN




RF

Network Traffic Classification Using ML & DL approaches


Precision Recall



CNN



CNN_LSTM




CNN_GRU

Summary and Future Work


Value of Temporal Features

- **Accuracy uplift**
- CNN + LSTM: +3.7% vs. static 1D-CNN
- CNN + GRU: +6.8% vs. static 1D-CNN
- **Flows most helped** : Audio、browsing
- **Flows less helped**: Chat
- **Costs & limitations**
 - more training time
 - Increased model complexity & memory footprint



Conclusions

- **Key findings**
- Temporal models spent more resources, but it does work
- **Model pros & cons**
- **K-NN**: no training, slow inference
- **RF**: robust, static
- **1D-CNN**: high cost, complex structure, low performance
- **CNN-LSTM/GRU**: higher cost, higher performance



Limitations and future directions

- **Limitations**
 - Dataset size is relatively small
 - Fixed windowing scheme may lose fine-grained temporal variations
- **Future Directions**
 - Develop adaptive windowing strategies to better capture dynamic timing patterns
 - Extend to multi-class application classification (e.g., YouTube, Facebook, Skype)

Thank you

Liu, Yucheng

